



Guix-HPC

Gestion des environnements logiciels dans l'ADT Gordon

(et petit historique de nos tentatives de déploiement logiciel)

Thursday March 7th, 2019

Outline

Context

Why providing support in a solver software stack?

Short review of earlier tentative deployments (CMake, Spack)

Preliminary results in ADT Gordon

Outline

Context

Why providing support in a solver software stack?

Short review of earlier tentative deployments (CMake, Spack)

Preliminary results in ADT Gordon

MORSE

Matrices Over Runtime Systems @ Exascale



University of Colorado
Denver



Linear algebra

$$\mathbf{A}\mathbf{X} = \mathbf{B}$$


Sequential-Task-Flow

```
for (j = 0; j < N; j++)  
  Task (A[j]);
```

Direct Acyclic Graph



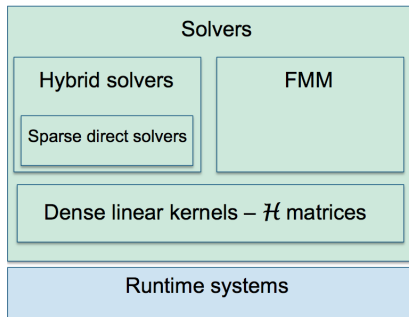
Runtime systems



Heterogeneous
platforms



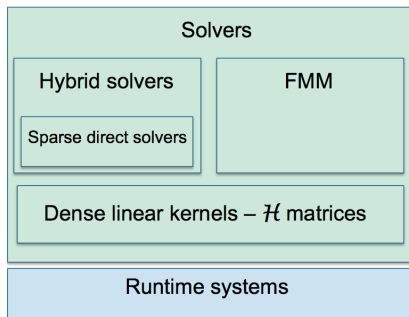
Numerical composition of solvers and runtime supports



Chameleon: dense linear solver

- Supported runtimes: **Quark** and **StarPU**, (**PaRSEC** soon)
- Ability to use cluster of heterogeneous nodes:
 - ▶ **MPI+threads**, CPUs (**BLAS/LAPACK**)+GPUs (**cuBLAS/MAGMA**)


Numerical composition of solvers and runtime supports



PaStiX: sparse linear solver

- LL^T , LDL^T , and LU , with static pivoting, supernodal approach
- Native version: $MPI+threads$
- Versions with runtimes: on top of $PaRSEC$ or $StarPU$

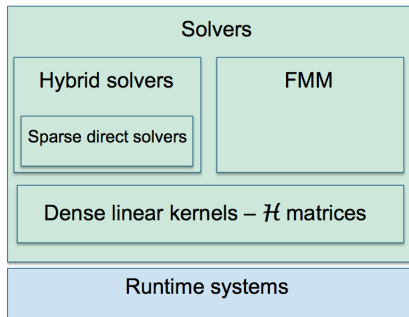
Numerical composition of solvers and runtime supports



MaPHyS: hybrid direct/iterative sparse linear solver

- Solves $\mathbf{Ax} = \mathbf{b}$, where \mathbf{A} is a square non singular general matrix
- Native version: **MPI**+**PaStiX**/**MUMPS**+**BLAS**/**LAPACK**
- Do not support runtimes for now, work in progress

Numerical composition of solvers and runtime supports




ScalFMM: scalable fast multipole methods

- Simulate N-body interactions using the Fast Multipole Method based on interpolation (Chebyshev or Lagrange)
- Native version: **MPI+OpenMP+BLAS+FFTW**
- Runtimes version: **StarPU, OpenMP4** → **StarPU** (ADT K'STAR)


Solver on top of runtime system - example: Chameleon

Chameleon : dense linear algebra tile algorithms (STF) on top of runtime systems

1) Tile matrix layout



3) Runtime systems



- QUARK: CPUs (pthread)
- StarPU: CPUs, CUDA, MPI

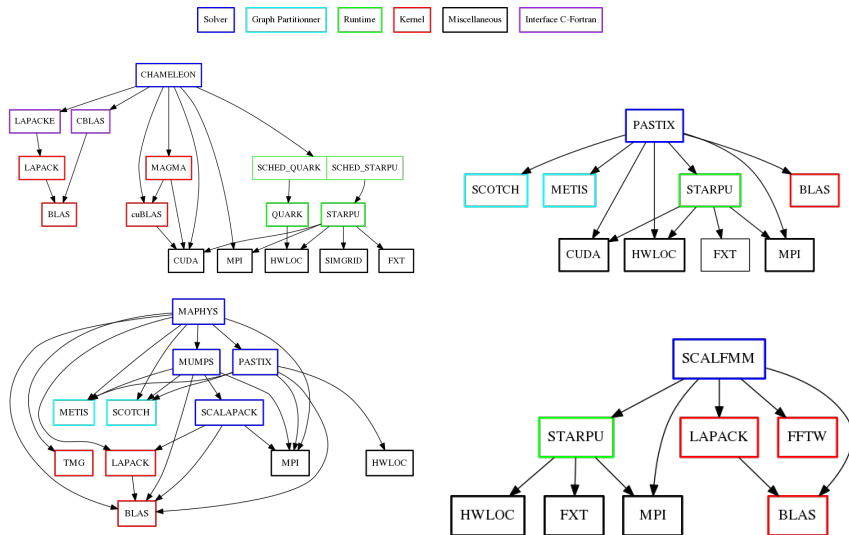
2) Algorithms

- STF Algorithms (**PLASMA**)
 - ▶ BLAS 3, some BLAS 2
 - ▶ LAPACK One-Sided, Norms
- Generic task submission interface
 - ▶ different runtimes can be plugged

4) Optimized kernels

- BLAS, LAPACK
 - ▶ Intel MKL, Netlib blas/Lapack, ...
- CUDA/cuBLAS, MAGMA
- Misc: Hwloc, FxT, SimGrid, EZTrace, ...

How to deploy complex HPC software stacks?



Outline

Context

Why providing support in a solver software stack?

Short review of earlier tentative deployments (CMake, Spack)

Preliminary results in ADT Gordon

Common situations

You spend **your life** on it: install the stack by yourself!

- specialists on all the stack are rare
- problems of compatibility between versions
- takes a lot of time

You want to **discover** a new solver

You want to contribute to a **particular feature**
(scheduling, runtime, ...)

- ask someone else to do it
- use pre-installed versions (binary packages, modules)
 - ▶ problem: only a couple of versions exist

Wish list: example of two profiles

1. Top-level users want the best version:
 - a default build with best options to get performances regarding the platform
2. A specialist wants to have the lead:
 - on the components he operates on
 - ▶ flexibility to set his version
 - but may not care about many dependencies
 - ▶ automatic choice of best options

Requirements

- A simple process to install a default version
- A flexible way to choose build variants
 - ▶ choose compiler, software versions
 - ▶ enable components, e.g. MPI : YES/NO
 - ▶ build options, e.g. `--enable-debug`
- Be able to install it on a remote cluster
 - ▶ no root permissions
 - ▶ no internet access (not necessarily)

Existing toolboxes

- PETSc
 - ▶ scientific library for solving PDEs in parallel MPI+Threads
 - ▶ wrappers to external solvers (partitionners, linear algebra, ...)
 - ▶ custom python scripts to activate packages
 - ▶ detection mode or download+install a web release, great !
 - ▶ detection problems, fixed versions to download
 - Trilinos
 - ▶ similar to PETSc, maybe even broader scope
 - ▶ embed one precise version of solvers
 - ▶ no tool to install missing third party libraries
- ⇒ **no competitive tool to install dependencies**

Outline

Context

Why providing support in a solver software stack?

Short review of earlier tentative deployments (CMake, Spack)

Preliminary results in ADT Gordon

Common process (C. Castagnède, F. Pruvost)

- use of CMake with similar options
- rely on the same detection of the system and libraries
 - ▶ recursive system of CMake Finds
 - ▶ if your application depends on [Chameleon](#), in CMake:

```
find_package(CHAMELEON COMPONENTS STARPU MPI CUDA MAGMA FXT)
```

List of available `find_package` in Morse:

solvers	chameleon, magma, mumps, pastix, plasma, scalapack
runtimes	quark, parsec, starpu
kernels	(c)blas, lapack(e), fftw
misc	(par)metis, (pt)scotch, hwloc, fxt, eztrace

State of the art: tool to distribute the software stack

- we do not want to reinvent the wheel *i.e.* use an existing solution:
 - ▶ Dpkg, 0install, Gub, Guix/Nix, Easybuild, ...
- classical package managers cannot meet our requirements
 - ▶ no root permissions, build variants easy to give, a mode to handle non open software (**Intel MKL**, **nvidia CUDA**)
- **Spack** a custom tool to install HPC libraries will be used



Spack

<http://scalability-llnl.github.io/spack/>

- Python 2.7: no install needed, ready to be used

```
$ git clone https://github.com/scalability-llnl/spack.git  
$ ./spack/bin/spack install gcc
```

- Easy way to set build variants, examples:

```
$ spack install openmpi %gcc@4.9.2  
$ spack install netlib-lapack +shared  
$ spack install parpack ^netlib-lapack ^openmpi@1.10.0
```

- Handle modulefiles, mirrors to work on clusters

```
$ spack load mpi  
$ spack mirror create openmpi mpich hwloc netlib-blas  
$ spack mirror add
```

Integration of our solvers (fork), F. Pruvost

- Build variants examples:

```
$ spack install maphys ~examples +mumps  
$ spack install pastix +starpu ^starpu@1.1.2 ^mkl-blas  
$ spack install starpu@svn-1.2 +debug +cuda +mpi +fxt +examples
```

Software integrated within **Spack** (fork)

Dense linear solvers

Chameleon, MAGMA, PLASMA,
ScaLAPACK

Sparse linear solvers

HIPS, MaPHyS, MUMPS, PaStiX,
qr_mumps, SuiteSparse

Fast Multipole Method solvers

ScalFMM

Runtime systems

QUARK, ParSEC, StarPU

Kernels

(C)BLAS (MKL, Netlib, OpenBlas),
LAPACK(E), FFTW

Miscellaneous

(Par)METIS, (PT-)SCOTCH, hwloc, MPI,
CUDA, SimGrid, EZTrace

Spack strengths and weaknesses

See Spack's slides.

Outline

Context

Why providing support in a solver software stack?

Short review of earlier tentative deployments (CMake, Spack)

Preliminary results in ADT Gordon

Preliminary results in ADT Gordon (A. Guilbaud, F. P.)

Current status

- ▶ Chameleon + StarPU + Open MPI
- ▶ Intel MKL
- ▶ Cuda (not much tested yet)
- ▶ Laptop / Plafrim

Preliminary results in ADT Gordon (A. Guilbaud, F. P.)

Current status

- ▶ Chameleon + StarPU + Open MPI
- ▶ Intel MKL
- ▶ Cuda (not much tested yet)
- ▶ Laptop / Plafrim
- ▶ Pastix, Maphys

Preliminary results in ADT Gordon (A. Guilbaud, F. P.)

Current status

- ▶ Chameleon + StarPU + Open MPI
- ▶ Intel MKL
- ▶ Cuda (not much tested yet)
- ▶ Laptop / Plafirm
- ▶ Pastix, Maphys

Next steps

- ▶ Further test Cuda, integrate Intel compilers
- ▶ Test deployment of the stack with **NewMadeleine** on **Plafirm**
- ▶ Further explore **variants** (“-with-source, -with-input, -with-branch, ...”)
- ▶ Handling **private** projects (FMR, Diodon)?

Some challenges

Make it a tool of choice

- ▶ for **end-users**
- ▶ for **contributors**
- ▶ for **core developers?**
- ▶ for **schools?**

The path is still long towards reproducibility

- ▶ Reproduce experiments
- ▶ Inception: simulation of the simulation (work of L. Stanisis)